


CS2109S Midterms

AY24/25 SEM 1
github/SelwynAng

1 Introduction to AI

1.1 Intelligent Agents



• PEAS

1. **Performance Measure:** Best for whom, what are we optimizing, what information is available, any unintended effects, what are the costs
 2. **Environment:** Refer to Environment section
 3. **Actuators:** Allow intelligent agent to take actions or affect its environment
 4. **Sensors:** Allow intelligent agent to perceive information about its environment
- **Agent Function:** Maps from percept histories to actions, refer to Agent Function section

1.2 Task Environment

1. Fully Observable VS Partially Observable

- Fully Observable: Agent has complete & accurate info about env state at all times (Eg. Chess)
- Partially Observable: Agent has access to incomplete, uncertain or noisy info about env state (Eg. Self-driving cars)

2. Deterministic VS Stochastic VS Strategic

- Deterministic: Next env state is completely determined by current state & agent action | Outcome is fully predictable (Eg. Sudoku)
- Stochastic: Next env state is not completely determined by current state & agent action | Outcome is uncertain (Eg. Self-driving car)
- Strategic: Env is deterministic, but outcomes depend on other agents' actions, requiring agent to consider strategies & behaviors of others (Eg. Chess)

3. Episodic VS Sequential

- Episodic: Agent actions are divided into discrete periods, each episode is independent of one another, agent makes decisions based on current episode (Eg. Classification task)
- Sequential: Agent actions are inter-dependent, each action affects future states & decisions, agent considers sequence of actions over time (Eg. Chess)

4. Static VS Dynamic

- Static: Env state does not change while agent is deliberating
- Dynamic: Env state changes over time even when agent is deliberating
- Semi-dynamic: Env state does not change, but agent's performance score does

5. Discrete VS Continuous

- Discrete: Finite # of distinct, clearly defined percepts & actions
- Continuous: Infinite # of percepts & actions

6. Single Agent VS Multi Agent

- Single Agent: Agent operating by itself in an env
- Multi Agent: Multiple agents in an env

1.3 Agent Structures

Note: Agent is completely specified by Agent Function mapping percept sequences to actions

1. **Simple Reflex Agent:** Operates based on a set of predefined rules or conditions → Reacts to current state of env with a corresponding action → Does not have memory of past states or actions & does not consider future consequences
 2. **Model-based Reflex Agent:** Extends simple reflex agent by maintaining internal model of world → Allows agent to keep track of current env state & handle situations where env state is partially observable or changes over time
 3. **Goal-based Agent:** Operates with specific goals in mind → Selects actions based on ability to achieve these goals → Considers future & plans its actions to achieve desired end state | Uses goal representation & perform search and planning
 4. **Utility-based Agent:** Extends goal based agent by considering not just whether goals are achieved, but how well they are achieved → Assigns utility value to different states & chooses actions that maximize overall utility
 5. **Learning Agent:** Improves performance over time by learning from its experiences | Can be reflex, model, goal & utility based
- **Exploitation:** Maximize expected utility according to current knowledge about world
 - **Exploration:** Trying to learn more about the world

2 Solving Problems by Searching

2.1 Designing an Agent

- **Assumptions:** Goal-based agent | Env is fully observable, deterministic, static, discrete
- **Problem-solving Agent:** Agent that plans ahead (considers a seq. of actions that form a path to a goal state), undertakes SEARCH process
- **Steps:**
 1. **Goal Formulation:** (What do we want?)

2. **Problem Formulation:** (How the world works?) → States (state space), Initial State(initial state of agent), Goal State/Test (goal state of agent), Actions (things that agent can do in a given state), Transition Model (specifies outcome of an action to a given state & how it leads to new states), Action Cost Function (cost of performing an action)
 3. **Search:** (How to achieve it?) → Path (seq. of actions), Solution (path to a goal)
 4. **Execute**
- **Representation Invariant:** A condition that must be true over all valid concrete representations of a class

2.2 Search Algorithms (Introduction)

- **Search Algorithm:** Takes in search problem (input), returns solution/failure (output) | Defined by Order of Expansion (FRONTIER)
- **Evaluation Criteria:**
 1. Time Complexity: # of nodes generated/expanded
 2. Space Complexity: Max # of nodes in memory
 3. Completeness: Does it return solution if it exists?
 4. Optimality: Does it always find least cost solution?

Tree Search:

create frontier

insert initial state to frontier

while frontier is not empty:

state = frontier.pop()

if state is goal: return solution

for action in actions(state):

next state = transition(state, action)

frontier.add(next state)

return failure

Graph Search:

create frontier

create visited

insert initial state to frontier

while frontier is not empty:

state = frontier.pop()

if state is goal: return solution

if state in visited: continue

visited.add(state)

for action in actions(state):

next state = transition(state, action)

frontier.add(next state)

return failure

• Checking of Goal State:

- New state is checked for goal state before new states are PUSHED to frontier → Expand less states, may skip states with less cost
- State is checked for goal state after state is POPPED from frontier → Expand more states, will not skip states with less cost

2.3 Search Algorithms (Uninformed Search)

- **Key Idea:** Search Algo is given no clue about how close a state is to the goal | Can be Tree or Graph Search
- **BFS:** Queue Frontier | Time Complexity: $O(b^d)$ = $1 + b + b^2 + \dots + b^d$, where b is branching factor, d is depth of optimal solution | Space Complexity: $O(b^d)$ when expanded until last child in worst case | Completeness: Complete if b is finite | Optimality: Optimal if step cost is same everywhere
- **UCS:** Priority Queue (path cost) Frontier, where path cost == cost from root to a state | Time Complexity: $O(b^{C^*}/\epsilon)$, where C^* is cost of optimal solution, ϵ is minimum edge cost → C^*/ϵ is est. depth of optimal solution in worst case | Completeness: Complete if $\epsilon > 0$ and C^* is finite (if $\epsilon = 0$, zero cost cycle may occur) | Optimality: Optimal if $\epsilon > 0$
- Note: BFS is special case of UCS where step cost == 1 for every edge
- **DFS:** Stack Frontier | Time Complexity: $O(b^m)$ where b is branching factor, m is max depth | Space Complexity: $O(bm)$ as only 1 path is expanded at one time | Completeness: Not complete (when depth is infinite or can go back or forth) | Optimality: Not optimal (there can be paths with less cost not explored yet)

- **DLS (Depth Limited Search):** Limit the search depth to l where $l < m$, backtrack once depth limit is reached | Time Complexity: $O(b^l)$ | Space Complexity: $O(b \cdot l)$ | Completeness: Not complete when soln lies deeper l | Optimality: Not optimal when soln lies deeper than l
- Note: We dk the depth of solution, which is a downside
- **IDS (Iterative Deepening Search):** Do DLS with max depth of $0, \dots, \infty$ → return soln if found, otherwise increase depth | Time Complexity: $O(b^d)$, Overhead = $(n_{IDS} - n_{DLS})/n_{DLS}$ | Space Complexity: $O(b \cdot d)$ | Completeness: Complete | Optimality: Optimal if step cost is same everywhere

- Note: IDS is not always faster than DFS → Consider state space s.t. each state have only single successor & goal node is at depth n → IDS will run in $O(n^2)$, DFS will run in $O(n)$
- **Backward Search:** Search from goal
- **Bidirectional Search:** Combine forward search & backward search, stop when 2 searches meet | Time Complexity: $2 * O(b^{d/2}) < O(b^d)$

2.4 Search Algorithms (Informed Search)

- **Key Idea:** Search Algo has a clue on how close a state is to the goal
- **Best First Search:** Priority Queue ($f(n)$) Frontier, where $f(n)$ estimates the goodness of a state (Node with lowest $f(n)$ is selected first to be expanded) | $f(n)$ can be purely heuristic (estimated cost from n to goal) or a combi of path cost & heuristic
- **Greedy Best First Search:** Priority Queue ($f(n) = h(n)$) Frontier, where $h(n)$ is heuristic function that est. cost from n to goal (Expands node that seems closest to goal according to $h(n)$ without considering path cost so far) | Time Complexity: $O(b^m)$ | Space Complexity: $O(b^m)$ | Completeness: Not complete since GBFS might keep expanding nodes based on $h(n)$ without ever finding goal | Optimality: Not optimal since GBFS selects nodes based on $h(n)$ without considering path cost
- **A* Search:** Priority Queue ($f(n) = g(n) + h(n)$) where $g(n)$ is cost so far to reach n | Time Complexity: $O(b^m)$ | Space Complexity: $O(b^m)$ | Completeness: Complete | Optimality: Optimal

- If $h(n)$ is admissible → A^* using Tree search is optimal
- If $h(n)$ is consistent → A^* using Graph search is optimal
- Note: UCS is special case of A^* search where $h(n) = 0$

2.5 Heuristics

- Estimate cost from node n to goal
- **Admissible Heuristics:** For every node n , $h(n) \leq h^*(n)$, where $h^*(n)$ is true cost to reach goal state from n (Never over-estimate)
- **Consistent Heuristics:** For every node n , every successor n' generated by action a , $h(n) \leq c(n, a, n') + h(n')$ and $h(G) = 0$ (Proof $h(n) - h(n') \leq c(n, a, n')$)
Note: If $h(n)$ is consistent, $f(n') \geq f(n) \rightarrow f(n)$ is non-decreasing along any path → Nodes are expanded in order of increasing f cost
- **Dominance:** If $h_2(n) \geq h_1(n)$ for all $n \rightarrow h_2$ dominates h_1 | If h_2 is admissible → h_2 is better for search
- **Creating Admissible Heuristics:**
 - Problem with fewer restrictions on actions is called a relaxed problem
 - Cost of an optimal soln to a relaxed problem is an admissible h for original problem

3 Local Search & Adversarial Search

3.1 Local Search

- **Assumptions:** Agent is a Goal/Utility-based agent, Env has a very large state space
- **Informed & Uninformed Search VS Local Search**
 1. **UCS:** Low to moderate state space | Optimal or no soln | Search path is usually the soln
 2. **LS:** Very large state space | Good enuf soln is preferable rather than no soln | State is the soln (don't care about search path)
- **Local Search Overview:**
 - **Basic Idea:** Start somewhere in state space, move towards a better spot
 - **Problem Formulation:** States(state space), Initial State(initial state of agent), Goal test (optional, coz we actually dk the goal state, rely on eval function instead), Successor Function (possible states from a state), Evaluation Function (Output value/goodness of a state)
- **Hill Climbing Algorithm**
current = initial state

while True:

neighbor = a highest-valued successor of current

if value(neighbor) <= value(current):

return current

current = neighbor

- Known as Greedy Local Search (pick best amongst neighbors, repeat)
- Best Soln: State space where eval. function has a max value (global max)
- Disadvantages: Cannot reach global max if it enters local max, plateau | Sensitive to choice of initial state, poor initial state may result in poor final state (Can overcome with random restarts, walks)

• Simulated Annealing

current = initial state

if a large positive value

while T > 0:

next = a randomly selected successor of current

if value(next) > value(current): current = next

else with probability P(current, next, T): current = next


decrease T

return current

- P(current, next, T) = $e^{(value(next) - value(current))/T}$
- More exploration of bad states is allowed when T is high, more exploitation is done when T is low → basically choosing worse successor may lead to a better max
- **Theorem:** If T decreases slowly enough, SA will find global optimum with high probability


3.2 Adversarial Search

- **Assumptions:** Agent is Utility-based | Env is a game (game cannot be single player, partially observable, stochastic, but must be fully observable, deterministic, discrete, terminal states exist, 2 players, zero-sum, turn taking)
- **Minimax Algorithm:**



- **Intuition:** MAX wins when utility is high, MIN wins when utility is low | Assign utility values to all terminal states & start tracing from terminal states → Eventually, all states will have utility values, starting player can choose a state that will max/min his utility
- **Analysis:** Completeness: Complete if tree is finite | Time Complexity: $O(b^m)$ | Space Complexity: $O(bm)$ depth first exploration | Optimality: Optimal against optimal opponent

• Alpha-beta Pruning




- **Definitions:** α is best explored option to the root for MAX player (Highest value for MAX) | β is best explored option along path to the root for MIN player (Lowest value for MIN)
- **Procedure:** 1. Assign $\alpha = -\infty$, $\beta = \infty$ for root 2. Propagate values down to the terminal node 3. Update α value at MAX node, β value at MIN node 4. Propagate values up 5. Prune branches of nodes where $\alpha \geq \beta$
- **Minimax with Cutoff**
 - Instead of calling *is.terminal*, call *is.cutoff* which returns TRUE if (1): State is terminal or (2): Cut-off is reached
 - Instead of using *utility*, call *eval* which is an eval. function that returns (1): Utility for terminal states or (2): Heuristic value for non-terminal states

4 Introduction to ML & Decision Trees

4.1 Introduction to ML

- **Definitions:** Computer program is said to learn from experience E w.r.t. some class of tasks T & performance measure P , if its performance at tasks in T , as measured by P , improves with experience E
- **Types of Feedback:**
 1. **Supervised Learning:** Involves training a model on a labeled dataset, where input data is paired with correct output → Model learns to map inputs to outputs based on this labeled data, allowing it to make predictions on new data
 - Regression: Predict continuous input
 - Classification: Predict discrete input
 2. **Unsupervised Learning:** Deals with dataset that do not have labeled outputs → Goal is to identify patterns & structures within data
 3. **Reinforcement Learning:** Agent learns to make decisions by interacting with an environment → Agent receives feedback in the form of rewards or penalties based on its actions → Learns optimal behaviors over time
- **Formal Definitions:**



4.2 Performance Measure

- **Regression:**

For a set of N examples $\{(x_1, y_1), \dots, (x_N, y_N)\}$ we can compute the average (mean) squared error as follows.

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

Where $\hat{y}_i = h(x_i)$ and $y_i = f(x_i)$.

- For a set of N examples $\{(x_1, y_1), \dots, (x_N, y_N)\}$ we can compute the average (mean) absolute error as follows.

$$MAE = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$$

Where $\hat{y}_i = h(x_i)$ and $y_i = f(x_i)$.

• Classification:

Classification is correct when the prediction $\hat{y} = y$ (true label).

- For a set of N examples $\{(x_1, y_1), \dots, (x_N, y_N)\}$ we can compute the average correctness (accuracy) as follows.

$$Accuracy = \frac{1}{N} \sum_{i=1}^N \mathbb{1}_{\hat{y}_i = y_i}$$

Where $\hat{y}_i = h(x_i)$ and $y_i = f(x_i)$.

		Actual Label	
		Cancer	Benign
Predicted Label	Cancer	2 True Positive	1 False Positive
	Benign	3 False Negative	4 True Negative

- **Accuracy:** $\frac{TP + TN}{TP + FN + FP + TN}$
- **Precision:** $\frac{TP}{TP + FP}$ (How many selected items are relevant, maximise if FP is costly)
- **Recall:** $\frac{TP}{TP + FN}$ (How many relevant items are selected, maximise if FN is dangerous)

• **F1 Score:** $\frac{2}{\frac{1}{precision} + \frac{1}{recall}}$

4.3 Decision Trees

• Traits of Decision Trees:

- Decision Trees can express any function of input attributes
- Consistent Decision Tree for any training set, but probably will not generalize to new examples
- # of distinct decision trees with n boolean attributes = 2^{2^n}

• Decision Tree Learning Algorithm

def DTL(examples, attributes, default):

if examples is empty: return default

if examples have the same classification: return classification

if attributes is empty: return mode(examples)

best = choose_attribute(attributes, examples)

tree = a new decision tree with root best

for each value v_i of best:

examples_i = {rows in examples with best = v_i }

subtree = DTL(examples_i, attributes - best, mode(examples_i))

add a branch to tree with label v_i and subtree subtree

- mode: Category with the highest number
- choose attribute: Chooses attribute with the highest information gain
- **Choosing an attribute:**
 - Ideally select an attribute that splits examples into "all positive" or "all negative"
 - Entropy (Measure of randomness): $I(P(v_1), \dots, P(v_n)) = -\sum_{i=1}^n P(v_i) \log_2 P(v_i)$, where for data set containing p positive & n negative examples, $I(\frac{p}{p+n}, \frac{n}{p+n}) = -\frac{p}{p+n} \log_2 \frac{p}{p+n} - \frac{n}{p+n} \log_2 \frac{n}{p+n}$
Note: $I(1, 0) = I(0, 1) = 0$, $I(0.5, 0.5) = 1$
 - Information Gain (Entropy of curr. node - Total Entropy of children nodes): $IG(A) = I(\frac{p}{p+n}, \frac{n}{p+n}) - remainder(A)$
remainder(A) = $\sum_{i=1}^v \frac{p_i + n_i}{p+n} I(\frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i})$, where examples are split into v subsets by attribute A

- **Dealing with continuous valued attributes:** Define a discrete valued input attribute to partition values into discrete set of intervals
- **Dealing with missing values:** Assign most common value of attribute, assign probability to each value and sample, drop attribute, drop rows
- **Overfitting:** Decision Tree is perfect on training data, but worse on test data
- **Occam Razor:** Prefer short/simple hypothesis (long/complex hypothesis that fits data may be coincidence)
- **Pruning:** Prevents nodes from being split even when it fails to cleanly separate examples (Min samples leaf: Merge until leaf node is above min. samples number | Max depth: Merge until leaf nodes are at depth less than max depth)

5 Midterms PYP Pointers

- **Tree Search VS Graph Search:** Assuming the problem consists of discrete states (eg. pitcher problem), search space using tree search may be infinite (Eg. we keep filling up and emptying pitcher) → but using graph search will make search space finite (graph search will not revisit visited states)
- **Heuristics:**
 - Max of 2 admissible heuristics ($max(h_0, h_1)$) for a problem is also an admissible heuristic → $max(h_0, h_1)$ is also dominant over h_0 & h_1 since it takes max of both heuristics
 - Given h_0 is admissible for problem p_0 & h_1 is admissible for problem p_1 → h_0 & h_1 are admissible for the combined problem $p_0 + p_1$ → $max(h_0, h_1)$ is also admissible and dominant for $p_0 + p_1$
 - If heuristic is consistent, it must be admissible too!